

Lab Course Machine Learning

Exercise Sheet 6

December 15th 2022/24 Janfr 311226

Exercise 0: Dataset Preprocessing

```
In [459]: import numpy as np
import math
from sklearn.preprocessing import PolynomialFeatures
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from matplotlib import cm
from matplotlib.ticker import LinearLocator
from sympy import symbols, diff
import pandas as pd
import math
import warnings
import itertools
warnings.filterwarnings('ignore')
from sklearn import model_selection
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
from sklearn.linear_model import ElasticNet
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.svm import SVR
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from math import sqrt
from sklearn.linear_model import SGDRegressor
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import StratifiedShuffleSplit

In [460]: mu, sigma = 1, 0.05 # mean and standard deviation
X = np.random.normal(mu, sigma, size=(100, 1))

In [461]: mu, sigma = 1, 0.1 # mean and standard deviation
r = np.random.normal(mu, sigma, size=(100, 1))

In [462]: Y = 1.3*(X**2) + (4.8*X) + 8.9

In [9]: wq = pd.read_csv('winequality-red.csv', sep = ',')

In [10]: wq

Out[10]:
```

| | fixed acidity | volatile acidity | citric acid | residual sugar | chlorides | free sulfur dioxide | total sulfur dioxide | density | pH | sulphates | alcohol | quality |
|------|---------------|------------------|-------------|----------------|-----------|---------------------|----------------------|---------|----------|-----------|---------|---------|
| 0 | 7.4 | 0.700 | 0.00 | 1.9 | 0.076 | 11.0 | 34.0 | 0.99780 | 3.51 | 0.56 | 9.4 | 5 |
| 1 | 7.8 | 0.880 | 0.00 | 2.6 | 0.098 | 25.0 | 67.0 | 0.99680 | 3.20 | 0.68 | 9.8 | 5 |
| 2 | 7.8 | 0.760 | 0.04 | 2.3 | 0.092 | 15.0 | 54.0 | 0.99700 | 3.26 | 0.65 | 9.8 | 6 |
| 3 | 11.2 | 0.280 | 0.56 | 1.9 | 0.075 | 17.0 | 60.0 | 0.99800 | 3.16 | 0.58 | 9.8 | 6 |
| 4 | 10.24 | 0.600 | 0.00 | 1.9 | 0.076 | 11.0 | 34.0 | 0.99540 | 3.51 | 0.56 | 9.4 | 5 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1594 | 6.2 | 0.600 | 0.08 | 2.0 | 0.090 | 32.0 | 44.0 | 0.99490 | 3.45 | 0.58 | 10.5 | 5 |
| 1595 | 5.9 | 0.550 | 0.10 | 2.2 | 0.062 | 39.0 | 51.0 | 0.99512 | 3.52 | 0.76 | 11.2 | 6 |
| 1596 | 6.28454 | 1.29660 | 0.06326 | 0.11438 | 0.181311 | 0.684967 | 0.314370 | 0.99603 | 0.481074 | 0.584594 | 6 | 6 |
| 1597 | 5.9 | 0.645 | 0.12 | 2.0 | 0.075 | 32.0 | 44.0 | 0.99547 | 3.57 | 0.71 | 10.2 | 5 |
| 1598 | 6.0 | 0.310 | 0.47 | 3.6 | 0.067 | 18.0 | 42.0 | 0.99549 | 3.39 | 0.66 | 11.0 | 6 |

1599 rows x 12 columns

```
In [11]: is_NaN = wq.isnull()
row_has_NaN = is_NaN.any(axis=1)
rows_with_NaN = wq[row_has_NaN]
len(rows_with_NaN)

Out[11]:
```

| | fixed acidity | volatile acidity | citric acid | residual sugar | chlorides | free sulfur dioxide | total sulfur dioxide | density | pH | sulphates | alcohol | quality |
|------|---------------|------------------|-------------|----------------|-----------|---------------------|----------------------|---------|----------|-----------|---------|---------|
| 0 | 7.4 | 0.700 | 0.00 | 1.9 | 0.076 | 11.0 | 34.0 | 0.99780 | 3.51 | 0.56 | 9.4 | 5 |
| 1 | 7.8 | 0.880 | 0.00 | 2.6 | 0.098 | 25.0 | 67.0 | 0.99680 | 3.20 | 0.68 | 9.8 | 5 |
| 2 | 7.8 | 0.760 | 0.04 | 2.3 | 0.092 | 15.0 | 54.0 | 0.99700 | 3.26 | 0.65 | 9.8 | 6 |
| 3 | 11.2 | 0.280 | 0.56 | 1.9 | 0.075 | 17.0 | 60.0 | 0.99800 | 3.16 | 0.58 | 9.8 | 6 |
| 4 | 10.24 | 0.600 | 0.00 | 1.9 | 0.076 | 11.0 | 34.0 | 0.99540 | 3.51 | 0.56 | 9.4 | 5 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1594 | 6.2 | 0.600 | 0.08 | 2.0 | 0.090 | 32.0 | 44.0 | 0.99490 | 3.45 | 0.58 | 10.5 | 5 |
| 1595 | 5.9 | 0.550 | 0.10 | 2.2 | 0.062 | 39.0 | 51.0 | 0.99512 | 3.52 | 0.76 | 11.2 | 6 |
| 1596 | 6.28454 | 1.29660 | 0.06326 | 0.11438 | 0.181311 | 0.684967 | 0.314370 | 0.99603 | 0.481074 | 0.584594 | 6 | 6 |
| 1597 | 5.9 | 0.645 | 0.12 | 2.0 | 0.075 | 32.0 | 44.0 | 0.99547 | 3.57 | 0.71 | 10.2 | 5 |
| 1598 | 6.0 | 0.310 | 0.47 | 3.6 | 0.067 | 18.0 | 42.0 | 0.99549 | 3.39 | 0.66 | 11.0 | 6 |

1599 rows x 12 columns

```
In [12]: dsy = wq.loc[:, wq.columns != 'quality']
wqt = (dsy-dsy.mean())/dsy.std()
wqt

Out[12]:
```

| | fixed acidity | volatile acidity | citric acid | residual sugar | chlorides | free sulfur dioxide | total sulfur dioxide | density | pH | sulphates | alcohol | quality |
|------|---------------|------------------|-------------|----------------|-----------|---------------------|----------------------|-----------|-----------|-----------|-----------|---------|
| 0 | -0.528194 | 0.961576 | -1.391037 | -0.453077 | -0.243630 | -0.466047 | -0.379014 | 0.558100 | -1.288240 | -0.579025 | -0.959946 | 5 |
| 1 | -0.298454 | 1.196827 | -1.391037 | 0.043403 | 0.223805 | 0.872365 | 0.624168 | 0.028252 | -0.719708 | 0.128910 | -0.584594 | 5 |
| 2 | -0.298454 | 1.196827 | -1.391037 | -0.169374 | 0.096323 | -0.083643 | 0.228975 | 0.134222 | -0.331073 | -0.048074 | -0.584594 | 6 |
| 3 | -0.298454 | 1.196827 | -1.391037 | -0.169374 | 0.096323 | -0.083643 | 0.228975 | 0.134222 | -0.331073 | -0.048074 | -0.584594 | 6 |
| 4 | -0.528194 | 0.961576 | -1.391037 | -0.453077 | -0.243630 | -0.466047 | -0.379014 | 0.558100 | -1.288240 | -0.579025 | -0.959946 | 5 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1594 | -1.217415 | 0.403103 | -0.980362 | -0.382151 | 0.053829 | 1.541571 | -0.075020 | -0.978459 | 0.899605 | -0.461036 | 0.072271 | 5 |
| 1595 | -1.389721 | 0.123866 | -0.877693 | -0.240300 | -0.541090 | 2.210777 | 0.137777 | -0.861893 | 1.353012 | 0.600867 | 0.291936 | 6 |
| 1596 | -1.159980 | -0.099523 | -0.723680 | -0.169374 | -0.243630 | 1.254769 | -0.196617 | -0.533387 | 0.705287 | 0.541872 | 0.541460 | 6 |
| 1597 | -1.389721 | 0.654416 | -0.750204 | -0.382151 | -0.264878 | 1.541571 | -0.075020 | -0.676446 | 1.676875 | 0.305894 | -0.209243 | 5 |
| 1598 | -1.332285 | -1.216469 | 1.021680 | 0.725659 | -0.434854 | 0.203159 | -0.135818 | -0.665849 | 0.510970 | 0.101921 | 0.541460 | 6 |

1599 rows x 12 columns

```
In [13]: wqt['quality'] = wq[['quality']]
wqt

Out[13]:
```

| | fixed acidity | volatile acidity | citric acid | residual sugar | chlorides | free sulfur dioxide | total sulfur dioxide | density | pH | sulphates | alcohol | quality |
|------|---------------|------------------|-------------|----------------|-----------|---------------------|----------------------|-----------|-----------|-----------|-----------|---------|
| 0 | -0.528194 | 0.961576 | -1.391037 | -0.453077 | -0.243630 | -0.466047 | -0.379014 | 0.558100 | -1.288240 | -0.579025 | -0.959946 | 5 |
| 1 | -0.298454 | 1.196827 | -1.391037 | 0.043403 | 0.223805 | 0.872365 | 0.624168 | 0.028252 | -0.719708 | 0.128910 | -0.584594 | 5 |
| 2 | -0.298454 | 1.196827 | -1.391037 | -0.169374 | 0.096323 | -0.083643 | 0.228975 | 0.134222 | -0.331073 | -0.048074 | -0.584594 | 6 |
| 3 | -0.298454 | 1.196827 | -1.391037 | -0.169374 | 0.096323 | -0.083643 | 0.228975 | 0.134222 | -0.331073 | -0.048074 | -0.584594 | 6 |
| 4 | -0.528194 | 0.961576 | -1.391037 | -0.453077 | -0.243630 | -0.466047 | -0.379014 | 0.558100 | -1.288240 | -0.579025 | -0.959946 | 5 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1594 | -1.217415 | 0.403103 | -0.980362 | -0.382151 | 0.053829 | 1.541571 | -0.075020 | -0.978459 | 0.899605 | -0.461036 | 0.072271 | 5 |
| 1595 | -1.389721 | 0.123866 | -0.877693 | -0.240300 | -0.541090 | 2.210777 | 0.137777 | -0.861893 | 1.353012 | 0.600867 | 0.291936 | 6 |
| 1596 | -1.159980 | -0.099523 | -0.723680 | -0.169374 | -0.243630 | 1.254769 | -0.196617 | -0.533387 | 0.705287 | 0.541872 | 0.541460 | 6 |
| 1597 | -1.389721 | 0.654416 | -0.750204 | -0.382151 | -0.264878 | 1.541571 | -0.075020 | -0.676446 | 1.676875 | 0.305894 | -0.209243 | 5 |
| 1598 | -1.332285 | -1.216469 | 1.021680 | 0.725659 | -0.434854 | 0.203159 | -0.135818 | -0.665849 | 0.510970 | 0.101921 | 0.541460 | 6 |

1599 rows x 12 columns

Exercise 1: Generalized Linear Models with Scikit Learn

Split your data into Train and Test Splits according to the 80%:20% ratio.

```
In [51]: wqt_train = wqt.loc[0:math.floor(len(wqt)*0.8)]
wqt_test = wqt.loc[math.floor(len(wqt)*0.8)+1:]
len(wqt_train)
len(wqt_test)

Out[51]:
```

| | fixed acidity | volatile acidity | citric acid | residual sugar | chlorides | free sulfur dioxide | total sulfur dioxide | density | pH | sulphates | alcohol | quality |
|------|---------------|------------------|-------------|----------------|-----------|---------------------|----------------------|-----------|-----------|-----------|-----------|---------|
| 1280 | -0.700500 | -0.378759 | -0.364349 | -0.453077 | -0.223283 | 1.159168 | 0.228975 | -0.607566 | 0.381425 | -0.107068 | -0.021567 | 6 |
| 1281 | -0.700500 | -0.378759 | -0.364349 | -0.453077 | -0.223283 | 1.159168 | 0.228975 | -0.607566 | 0.381425 | -0.107068 | -0.021567 | 6 |
| 1282 | -0.241019 | 1.324589 | -1.391037 | -0.114381 | -0.073654 | -0.055748 | -0.743808 | -0.294955 | 0.123335 | -0.126910 | 0.447623 | 6 |
| 1283 | 0.128462 | 0.570645 | 0.046326 | 0.114328 | 0.181311 | 0.672588 | 0.411372 | 0.664069 | -0.978798 | -0.461036 | -0.584594 | 6 |
| 1284 | -0.573935 | -0.602148 | -0.415683 | -0.169374 | -0.348666 | 0.203159 | -0.318215 | -0.105263 | 0.510970 | -0.579025 | 0.447623 | 5 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1594 | -1.217415 | 0.403103 | -0.980362 | -0.382151 | 0.053829 | 1.541571 | -0.075020 | -0.978459 | 0.899605 | -0.461036 | 0.072271 | 5 |
| 1595 | -1.389721 | 0.123866 | -0.877693 | -0.240300 | -0.541090 | 2.210777 | 0.137777 | -0.861893 | 1.353012 | 0.600867 | 0.291936 | 6 |
| 1596 | -1.159980 | -0.099523 | -0.723680 | -0.169374 | -0.243630 | 1.254769 | -0.196617 | -0.533387 | 0.705287 | 0.541872 | 0.541460 | 6 |
| 1597 | -1.389721 | 0.654416 | -0.750204 | -0.382151 | -0.264878 | 1.541571 | -0.075020 | -0.676446 | 1.676875 | 0.305894 | -0.209243 | 5 |
| 1598 | -1.332285 | -1.216469 | 1.021680 | 0.725659 | -0.434854 | 0.203159 | -0.135818 | -0.665849 | 0.510970 | 0.101921 | 0.541460 | 6 |

319 rows x 12 columns

For each model, pick three sets of hyperparameters and learn each model

```
In [479]: X_train = wqt_train.loc[:, wqt_train.columns != 'quality'].to_numpy()
Y_train = wqt_train['quality'].to_numpy()

X_test = wqt_test.loc[:, wqt_test.columns != 'quality'].to_numpy()
Y_test = wqt_test['quality'].to_numpy()

In [480]: n = 3
lamdaList = np.random.uniform(np.exp(-3), np.exp(-4), n)
numberIterations = [250, 500, 1000]
for i in range(n):
    hyperParamConfigs.append([numberIterations[i], lamdaList[i]])

In [481]: hyperParamConfigs

Out[481]:
```

| | numberIterations | lamda |
|---|------------------|----------------------|
| 0 | 250 | 0.010442327632894041 |
| 1 | 500 | 0.010112096005841741 |
| 2 | 1000 | 0.017106943712182084 |

```
In [482]: Rmse = []
for i in range(n):
    ordinaryReg = SGDRegressor(alpha = lamda, max_iter=numberIterations[i])
    pred_train_ordinaryReg = ordinaryReg.predict(X_train)
    rmseTrain_ordinaryReg = np.sqrt(mean_squared_error(Y_train, pred_train_ordinaryReg))

    ridgeReg = Ridge(alpha = lamda, max_iter=numberIterations[i], penalty='l2')
    ridgeReg.fit(X_train, Y_train)
    pred_train_ridgeReg = ridgeReg.predict(X_train)
    rmseTrain_ridgeReg = np.sqrt(mean_squared_error(Y_train, pred_train_ridgeReg))

    lassoReg = Lasso(alpha = lamda, max_iter=numberIterations[i], penalty='l1')
    lassoReg.fit(X_train, Y_train)
    pred_train_lassoReg = lassoReg.predict(X_train)
    rmseTrain_lassoReg = np.sqrt(mean_squared_error(Y_train, pred_train_lassoReg))

    pred_test_ordinaryReg = ordinaryReg.predict(X_test)
    rmseTest_ordinaryReg = np.sqrt(mean_squared_error(Y_test, pred_test_ordinaryReg))

    pred_test_ridgeReg = ridgeReg.predict(X_test)
    rmseTest_ridgeReg = np.sqrt(mean_squared_error(Y_test, pred_test_ridgeReg))

    pred_test_lassoReg = lassoReg.predict(X_test)
    rmseTest_lassoReg = np.sqrt(mean_squared_error(Y_test, pred_test_lassoReg))

    RMSE.append([i, lamda, rmseTrain_ordinaryReg, rmseTest_ordinaryReg, rmseTrain_ridgeReg, rmseTest_ridgeReg, rmseTrain_lassoReg, rmseTest_lassoReg])

In [483]: Rmse = np.array(RMSE)

In [484]: np.array(hyperParamConfigs)[i,:]
```

array([0.01044232, 0.0101121, 0.01710694])

```
In [508]: fig = plt.figure(figsize=(14,10))

ax = fig.add_subplot(1, 1, 1, projection='3d')
ax.view_init(elev=10, azim=80)
ax.plot(Rmse[:,0], Rmse[:,1], Rmse[:,4], label='rmseTrain_ordinary')
ax.plot(Rmse[:,0], Rmse[:,1], Rmse[:,5], label='rmseTest_ordinary')
ax.set_xlabel('Iterations')
ax.set_ylabel('Lambda')
ax.set_zlabel('RMSE')

ax.set_title('Ordinary Least Squares')

ax = fig.add_subplot(1, 1, 1, projection='3d')
ax.plot(Rmse[:,0], Rmse[:,1], Rmse[:,2], label='rmseTrain_ridgeReg')
ax.plot(Rmse[:,0], Rmse[:,1], Rmse[:,3], label='rmseTest_ridgeReg')
ax.set_xlabel('Iterations')
ax.set_ylabel('Lambda')
ax.set_zlabel('RMSE')

ax.set_title('Ridge Regression')

ax = fig.add_subplot(1, 1, 1, projection='3d')
ax.plot(Rmse[:,0], Rmse[:,1], Rmse[:,6], label='rmseTrain_lassoReg')
ax.plot(Rmse[:,0], Rmse[:,1], Rmse[:,7], label='rmseTest_lassoReg')
ax.set_xlabel('Iterations')
ax.set_ylabel('Lambda')
ax.set_zlabel('RMSE')

ax.set_title('LASSO')

plt.show()
```

As it is clear from graph that train error is less than test error

3. Now tune the hyperparameters using scikit learn GridSearchCV

```
In [639]: def get_best_params(train_X, train_Y, alphaRange, iterationRange, penalty):
    param_grid = {'alpha':alphaRange, 'max_iter':iterationRange}
    cv = StratifiedShuffleSplit(n_splits=10, test_size=0.2, random_state=42)
    for param_grid calling GridSearchCV with penalty (or specifying model) and diff hyper params in grid
    grid = GridSearchCV(estimator=SGDRegressor(penalty=penalty), param_grid=param_grid, cv=cv, return_train_score=True, n_jobs=-1)
    grid.fit(train_X, train_Y)

    print("The best parameters are %s with a score of %0.2f" % (grid.best_params_, grid.best_score_))

    scores = []
    for i,j,k in zip(grid.cv_results_['params'], grid.cv_results_['mean_test_score'], grid.cv_results_['mean_train_score']):
        scores.append((j,k,i['alpha'],i['max_iter']))
    no_iteration_best = grid.best_params_['max_iter']
    alpha_best = grid.best_params_['alpha']

    sgd = SGDRegressor(alpha = alpha_best, max_iter=no_iteration_best, penalty=penalty)

    sgd.fit(train_X, train_Y)

    X_train = wqt_train.loc[:, wqt_train.columns != 'quality'].to_numpy()
    Y_train = wqt_train['quality'].to_numpy()

    X_test = wqt_test.loc[:, wqt_test.columns != 'quality'].to_numpy()
    Y_test = wqt_test['quality'].to_numpy()

Ordinary Least Squares

iterationRange = [1, 10, 100, 1000]
alphaRange = [0]
ordinarySGD, scores, no_iteration_best_ord, alpha_best_ord = get_best_params(X_train, Y_train, alphaRange, iterationRange)
fig = plt.figure(figsize=(14,10))
ax = fig.add_subplot(1, 1, 1, projection='3d')
ax.scatter(scores[:,2], scores[:,3], scores[:,0], label='test')
ax.scatter(scores[:,2], scores[:,3], scores[:,1], label='train')
ax.set_xlabel('Iterations')
ax.set_ylabel('Lambda')
ax.set_zlabel('Accuracy Score')
ax.legend()
plt.show()

Fitting 2 folds for each of 4 candidates, totalling 8 fits
The best parameters are ('alpha': 0, 'max_iter': 1000) with a score of 0.35
[Parallel(n_jobs=4)]: Using backend LokyBackend with 14 concurrent workers.
[Parallel(n_jobs=4)]: Batch computation too fast (0.010s). Setting batch_size=2.
[Parallel(n_jobs=4)]: Done 4 tasks | elapsed: 0.0s
[Parallel(n_jobs=4)]: Done 13 tasks | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=4)]: Done 18 out of 40 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=4)]: Done 23 out of 40 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=4)]: Done 28 out of 40 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=4)]: Done 33 out of 40 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=4)]: Done 38 out of 40 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=4)]: Done 40 out of 40 | elapsed: 0.0s finished

Ridge Regression

iterationRange = [1, 10, 100, 1000]
alphaRange = np.random.uniform(np.exp(-9), np.exp(-4), n)
ridgeSGD, scores, no_iteration_best_ri, alpha_best_ri = get_best_params(X_train, Y_train, alphaRange, iterationRange)
fig = plt.figure(figsize=(14,10))
ax = fig.add_subplot(1, 1, 1, projection='3d')
ax.scatter(scores[:,2], scores[:,3], scores[:,0], label='test')
ax.scatter(scores[:,2], scores[:,3], scores[:,1], label='train')
ax.set_xlabel('Iterations')
ax.set_ylabel('Lambda')
ax.set_zlabel('Accuracy Score')
ax.legend()
plt.show()

Fitting 2 folds for each of 20 candidates, totalling 40 fits
The best parameters are ('alpha': 0.000581469970294834, 'max_iter': 1000) with a score of 0.35
[Parallel(n_jobs=4)]: Using backend LokyBackend with 14 concurrent workers.
[Parallel(n_jobs=4)]: Batch computation too fast (0.010s). Setting batch_size=2.
[Parallel(n_jobs=4)]: Done 4 tasks | elapsed: 0.0s
[Parallel(n_jobs=4)]: Done 13 tasks | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=4)]: Done 18 out of 40 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=4)]: Done 23 out of 40 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=4)]: Done 28 out of 40 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=4)]: Done 33 out of 40 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=4)]: Done 38 out of 40 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=4)]: Done 40 out of 40 | elapsed: 0.0s finished

LASSO

iterationRange = [1, 10, 100, 1000]
alphaRange = np.random.uniform(np.exp(-9), np.exp(-4), n)
lassoSGD, scores, no_iteration_best_las, alpha_best_las = get_best_params(X_train, Y_train, alphaRange, iterationRange)
fig = plt.figure(figsize=(14,10))
ax = fig.add_subplot(1, 1, 1, projection='3d')
ax.scatter(scores[:,2], scores[:,3], scores[:,0], label='test')
ax.scatter(scores[:,2], scores[:,3], scores[:,1], label='train')
ax.set_xlabel('Iterations')
ax.set_ylabel('Lambda')
ax.set_zlabel('Accuracy Score')
ax.legend()
plt.show()

Fitting 2 folds for each of 20 candidates, totalling 40 fits
The best parameters are ('alpha': 0.0037581469970294834, 'max_iter': 1000) with a score of 0.35
[Parallel(n_jobs=4)]: Using backend LokyBackend with 14 concurrent workers.
[Parallel(n_jobs=4)]: Batch computation too fast (0.010s). Setting batch_size=2.
[Parallel(n_jobs=4)]: Done 4 tasks | elapsed: 0.0s
[Parallel(n_jobs=4)]: Done 13 tasks | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=4)]: Done 18 out of 40 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=4)]: Done 23 out of 40 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=4)]: Done 28 out of 40 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=4)]: Done 33 out of 40 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=4)]: Done 38 out of 40 | elapsed: 0.0s remaining: 0.0s
[Parallel
```