

Lab Course Machine Learning

Exercise Sheet 8


January 5th, 2022

Syed Wasif Murtaza Jafri-311226

Exercise 1: Optical Character Recognition via Neural Networks

```
In [23]: import numpy as np
import math
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from matplotlib import cm
from matplotlib.ticker import LinearLocator
plt.rcParams['figure.figsize'] = (10, 8)
from sympy import symbols, diff
import pandas as pd
import math
import warnings
import itertools
warnings.filterwarnings('ignore')
import sklearn
from sklearn.model_selection import train_test_split
from sklearn.datasets import load_digits
from sklearn.neural_network import MLPClassifier
from sklearn import model_selection
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
from sklearn.linear_model import ElasticNet
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.svm import SVR
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from math import sqrt
from sklearn.linear_model import SGDRegressor
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import StratifiedShuffleSplit
import torch.nn.functional as F
from torchvision import transforms
import os
import math
from torch.utils.data.sampler import SequentialSampler
import torch
from torch import nn
from torch.utils.data import DataLoader, Dataset
import cv2
import os
from sklearn.metrics import accuracy_score
```

```
In [7]: digits = load_digits(as_frame=False)
plt.gray()
plt.matshow(digits.images[110])
plt.show()
```



```
In [8]: X = digits.data

In [9]: Y = digits.target
Y = Y.reshape(len(Y),1)

In [10]: X_train, X_test, y_train, y_test = train_test_split(X, Y, stratify=Y, test_size=0.2, random_state=3116)

In [11]: clf = MLPClassifier(random_state=3116, max_iter=300).fit(X_train, y_train)

In [12]: clf.predict_proba(X_test[:1])

Out[12]: array([[2.54100615e-08, 2.62938165e-08, 2.13561410e-15, 4.73356379e-20,
9.99999868e-01, 3.41843016e-11, 1.40238884e-08, 6.61206535e-08,
6.16084626e-12, 1.16212870e-11]])

In [13]: clf.predict(X_test[:1])

Out[13]: array([4])

In [14]: clf.predict(X_test[:5, :])

Out[14]: array([4, 5, 2, 3, 3])

In [15]: clf.score(X_test, y_test)

Out[15]: 0.9666666666666667

In [16]: def get_best_params(train_x, train_y, hidden_layer_sizes_range, batch_size_range):
    param_grid = dict(hidden_layer_sizes=hidden_layer_sizes_range, batch_size=batch_size_range)

    cv = StratifiedShuffleSplit(n_splits=2, test_size=0.2, random_state=42) # stratifying sample for crossvalidation
    # for param grid calling GridSearchCV with penalty (for specifying model) and diff hyper params in grid
    grid = GridSearchCV(MLPClassifier(random_state=3116, max_iter=500), param_grid=param_grid, cv=cv, return_train_score=True)
    grid.fit(train_x, train_y)

    print("The best parameters are %s with a score of %0.2f" % (grid.best_params_, grid.best_score_))

    scores = []
    # to calculate test and train score for each setting in parameter and adding it to scores list
    for i, j, k in zip(grid.cv_results_['params'], grid.cv_results_['mean_test_score'], grid.cv_results_['mean_train_score']):
        scores.append([j, k, i['hidden_layer_sizes'], i['batch_size']])
    hidden_layer_sizes_best = grid.best_params_['hidden_layer_sizes']
    batch_sizebest = grid.best_params_['batch_size']

    clf = MLPClassifier(hidden_layer_sizes =hidden_layer_sizes_best, batch_size=batch_sizebest)

    return clf, np.array(scores), hidden_layer_sizes_best, batch_sizebest
```

```
In [33]: hidden_layer_sizes_range = [2, 16, 32, 64, 100]
batch_size_range = [16, 32, 64, 128, 264, 528]
clf, scores, hidden_layer_sizes_best, batch_sizebest = get_best_params(X_train, y_train, hidden_layer_sizes_range, batch_size_range)

Fitting 2 folds for each of 30 candidates, totalling 60 fits
[Parallel(n_jobs=14)]: Using backend LokyBackend with 14 concurrent workers.
[Parallel(n_jobs=14)]: Done 4 tasks | elapsed: 6.9s
[Parallel(n_jobs=14)]: Done 13 tasks | elapsed: 9.5s
[Parallel(n_jobs=14)]: Done 22 tasks | elapsed: 13.2s
[Parallel(n_jobs=14)]: Done 33 tasks | elapsed: 16.4s
[Parallel(n_jobs=14)]: Done 40 out of 60 | elapsed: 17.9s remaining: 8.9s
[Parallel(n_jobs=14)]: Done 47 out of 60 | elapsed: 19.7s remaining: 5.4s
[Parallel(n_jobs=14)]: Done 54 out of 60 | elapsed: 21.0s remaining: 2.2s
[Parallel(n_jobs=14)]: Done 60 out of 60 | elapsed: 21.7s finished
The best parameters are {'batch_size': 16, 'hidden_layer_sizes': 100} with a score of 0.98

In [34]: clf.fit(X_train, y_train)

Out[34]: MLPClassifier(batch_size=16, hidden_layer_sizes=100)

In [35]: y_pred = clf.predict(X_test)

In [36]: acc = accuracy_score(y_test, y_pred)

In [37]: print('Test Accuracy::', acc)

Test Accuracy:: 0.9666666666666667

In [38]: print("The best parameters are: batch_size = %i and hidden_layer_sizes = %i" % (batch_sizebest, hidden_layer_sizes_best))

The best parameters are: batch_size = 16 and hidden_layer_sizes = 100
```

Exercise 2: End-to-End Self-Driving via Convolutional Neural Networks

```
In [80]: device = "cuda" if torch.cuda.is_available() else "cpu"
class Dataset(Dataset):

    def __init__(self):
        # default directory where data is loaded
        self.filepath = 'driving_dataset/'
        self.filenames = os.listdir(self.filepath)
        self.xs = []
        self.ys = []
        with open("angles.txt") as f:
            for line in f:
                self.xs.append("driving_dataset/" + line.split()[0])
                self.ys.append(float(line.split()[1]) * 3.14159265 / 180)

    def __len__(self):
        return len(self.filenames)

    def __getitem__(self, index):
        filename = self.filenames[index]
        img = cv2.imread(self.filepath + filename)
        height, width, channels = img.shape
        # print(height, width, channels)
        # Resizing image to (32, 32)
        resized = cv2.resize(img, (66, 200), interpolation=cv2.INTER_AREA)
        transform_norm = transforms.Compose([
            transforms.ToTensor()
        ])
        # to Tensor
        tensorImg2 = transform_norm(resized)

        mean = tensorImg2.mean([1, 2])
        std = tensorImg2.std([1, 2])

        transform_norm = transforms.Compose([
            transforms.ToTensor(),
            transforms.Normalize(mean, std)
        ])
        norm_tensorImg2 = transform_norm(resized)

        return norm_tensorImg2.float(), torch.tensor([self.ys[index]])

class ConvNet(torch.nn.Module):
    def __init__(self):
        super().__init__()

        self.conv1 = nn.Conv2d(3, 24, 5, stride=(2, 2))
        self.conv2 = nn.Conv2d(24, 36, 5, stride=(2, 2))
        self.conv3 = nn.Conv2d(36, 48, 5, stride=(2, 2))
        self.conv4 = nn.Conv2d(48, 64, 3)
        self.conv5 = nn.Conv2d(64, 64, 3)
        self.fc1 = nn.Linear(64 * 18 * 1, 100)
        self.fc2 = nn.Linear(100, 50)
        self.fc3 = nn.Linear(50, 10)
        self.out = nn.Linear(10, 1)

    def forward(self, x):
        # input x size : torch.Size([60, 3, 200, 66])

        # 5 Convolution Layers
        x = F.relu(self.conv1(x)) # 1 CV 1 torch.Size([60, 24, 98, 31])
        x = F.relu(self.conv2(x)) # CV 2 torch.Size([60, 36, 47, 14])
        x = F.relu(self.conv3(x)) # CV 3 torch.Size([60, 36, 47, 14])
        x = F.relu(self.conv4(x)) # CV 4 torch.Size([60, 64, 20, 3])
        x = F.relu(self.conv5(x)) # CV 5 torch.Size([60, 64, 18, 1])

        x = x.view(-1, 64 * 18 * 1) # flatten torch.Size([60, 1152])

        # 3 Fully connected Layers
        x = F.relu(self.fc1(x)) # torch.Size([60, 100])
        x = F.relu(self.fc2(x)) # torch.Size([60, 50])
        x = F.relu(self.fc3(x)) # torch.Size([60, 10])

        # Output layer with activation function
        x = F.tanh(self.out(x)) # out torch.Size([60, 1])

        return x

In [154]: dataset = Dataset()

batch_size = 60
shuffle_dataset = False
random_seed = 31142
num_epochs = 5

# Creating data indices for training and validation splits:
dataset_size = len(dataset)

indices = list(range(dataset_size))
if shuffle_dataset:
    np.random.seed(random_seed)
    np.random.shuffle(indices)

testSplitRatio = 10000 / dataset_size
validationSplitRatio = 10000 / dataset_size
trainSplitRatio = 1 - testSplitRatio - validationSplitRatio

train_indices = indices[0:math.floor(dataset_size * trainSplitRatio)]
val_indices = indices[math.floor(dataset_size * trainSplitRatio):math.floor(dataset_size * (trainSplitRatio + validationSplitRatio))]
test_indices = indices[math.floor(dataset_size * (trainSplitRatio + validationSplitRatio)):]

train_sampler = SequentialSampler(train_indices)
valid_sampler = SequentialSampler(val_indices)
test_sampler = SequentialSampler(test_indices)

train_loader = torch.utils.data.DataLoader(dataset, batch_size=batch_size,
                                            sampler=train_sampler, shuffle=False)
validation_loader = torch.utils.data.DataLoader(dataset, batch_size=batch_size,
                                                sampler=valid_sampler, shuffle=False)
test_loader = torch.utils.data.DataLoader(dataset,
                                           sampler=test_sampler, shuffle=False)

10000

In [159]: n_total_steps = len(train_loader)
net = ConvNet().to(device)
optimizer = torch.optim.Adam(net.parameters(), lr=1e-3)
criterion = torch.nn.MSELoss()
for epoch in range(num_epochs):
    for i, sample_batched in enumerate(train_loader):
        optimizer.zero_grad()
        yhat = net(sample_batched[0].cuda())
        loss = criterion(yhat.squeeze(), sample_batched[1].squeeze()).cuda()
        loss.backward(retain_graph=True)
        optimizer.step()
        if (i+1) * 200 == 0:
            print(f'Epoch [(epoch+1)/(num_epochs)], Step [(i+1)/(n_total_steps)], Loss: {loss.item():.4f}')

    print('Finished Training')

Epoch [1/5], Step [200/427], Loss: 0.0612
Epoch [1/5], Step [400/427], Loss: 0.0335
Epoch [2/5], Step [200/427], Loss: 0.0679
Epoch [2/5], Step [400/427], Loss: 0.0315
Epoch [3/5], Step [200/427], Loss: 0.0732
Epoch [3/5], Step [400/427], Loss: 0.0284
Epoch [4/5], Step [200/427], Loss: 0.0737
Epoch [4/5], Step [400/427], Loss: 0.0277
Epoch [5/5], Step [200/427], Loss: 0.0742
Epoch [5/5], Step [400/427], Loss: 0.0271
Finished Training

In [161]: def MSELoss(yhat, y):
    return torch.mean((yhat - y) ** 2)
criterion = MSELoss
mse = 0
with torch.no_grad():
    for i, sample_batched in enumerate(test_loader):
        optimizer.zero_grad()
        yhat = net(sample_batched[0].cuda())
        loss = criterion(yhat.squeeze(), sample_batched[1].squeeze()).cuda()
        mse += loss
    rmse = torch.sqrt(mse)
    rmseAvg = rmse / len(test_loader)

rmseAvg

tensor(0.0077, device='cuda:0')
```

```
In [158]: rmseAvg.item()

Out [158]: 0.00759779242798686

In [160]: print('RMSE on test images', rmseAvg.item())

RMSE on test images 0.00759779242798686
```